

# Investigate mono repo

<https://issues.redhat.com/browse/RUN-2785>

## Goal

To enhance the release process, dev workflow and reduce the vendoring issues, one of the options suggested is to merge all the container repos (**polyrepos**) into one repo (**monorepo**).

This document discusses this goal and tries to find out the pros and cons for this approach as well as for the ways how to do this migration.

## Scope

Following repositories are in scope:

- <https://github.com/containers/storage>
- <https://github.com/containers/image>
- <https://github.com/containers/common>

## Pros and cons for monorepo

### Pros

#### 1. Atomic Changes Across Modules

We can make a single commit that updates code in common, and applies the change across storage, and image - all in one go, without coordination across repos or version bumps.

#### 2. Simplified Dependency Management

No need to version, tag, or publish common, image, and storage separately. No more replace directives in go.mod to test local changes across modules.

#### 3. Unified CI/CD

One pipeline can test the integration of all parts. Easier to ensure compatibility between storage, image, and shared packages. Shared tooling (e.g. linters, build scripts) can live in a single place.

#### 4. Better Code Navigation

For contributors and new devs: Easier to explore code and understand relationships between modules. No need to jump across repos and clone multiple things.

## 5. Simplified Project Coordination

All code lives in one place, making:

- Onboarding easier
- Internal refactors smoother
- Cross-team communication tighter
- Single issue tracker to check

## Cons

### 1. Potential for Code Bloat

The monorepo grows large, and can:

- Slow down operations like cloning or searching
- Increase build times and CI test times.

### 2. Harder Module Independence

Possibility to lose clear API boundaries that polyrepos naturally enforce. Risk of tighter coupling.

### 3. Migration cost

Migrating to a monorepo involves:

- Git history preservation
- Combining and deduplicating CI, build tools
- Refactoring import paths and go.mod files

### 4. Tooling at Scale

Monorepos require better tooling:

- Dependency-aware builds (e.g. bazel, mage, or advanced make)
- Selective testing (only run tests for affected components)
- Maybe go.work (Go 1.18+) or mono-module setups

### 4. Single issue tracker

Using a single repository also means using a single issue tracker. We would need to change our workflow to support that.

# Monorepo structure

## Open questions

- Do we want to treat image, storage and common as a single module, or do we want to treat them as three separate modules with separate releases, but living in a single monorepo?

## Option 1: Standard Go Layout

Unset

```
monorepo/
├── cmd/                # Entry points (main.go)
│   ├── image/
│   └── storage/
├── docs/              # Can be shared between projects.
├── hack/              # Can be shared between projects.
├── pkg/               # Reusable libraries (exported APIs)
│   ├── common/
│   ├── image/
│   └── storage/
├── internal/          # Private shared code
│   ├── common/
│   ├── image/
│   ├── storage/
│   └── drivers/       # Moved from /drivers.
├── tests/             # Tests for each project.
│   ├── common/
│   ├── image/
│   └── storage/
├── types/             # Shared structs, interfaces, enums, constants, ...
│   ├── common/
│   ├── image/
│   └── storage/
├── vendor/            # Can be shared between projects.
└── go.mod             # Single go module
```

### Pros:

- Familiar to most Go developers (community convention)
- Clean separation of concerns (binaries vs. public libraries vs internal libraries)
- Works well for tooling and CLIs
- go mod vendor works in a single step (single vendor directory)

- No need to manage versions between internal packages (if we want to mix them).

#### Cons:

- Not ideal if we want completely separate modules or versioning
- Shared go.mod can get large over time

### Option 2: Multi-Module Layout (With go.work)

It is a workspace file (go.work) that ties multiple go.mod projects together in a single local development environment.

```
Unset
monorepo/
├── image/                # Independent Go module
│   ├── ...
│   └── go.mod
├── storage/             # Independent Go module
│   ├── ...
│   └── go.mod
├── common/              # Independent Go module
│   ├── ...
│   └── go.mod
└── go.work              # Ties them together
```

#### Pros:

- Keeps module independence (e.g., separate releases possible), even inside the monorepo.
- Great for gradual migration from polyrepo (easy to migrate)
- Forces explicit APIs between parts (retain boundaries between projects)
- We can easily test and build single module in isolation.

#### Cons:

- More go.mod files = more boilerplate, but not more than what we have now.
- CI/CD and dependency management are more complex, but easier for developers than what we have now.

Example go.work:

```
Unset
go 1.22

use (
```

```
    ./image
    ./storage
    ./common
)
```

This tells Go:

- When we run go build, go test, go run, etc.,
- If a package inside image/ imports something from common/,
- Go will directly use the local source code from ./common/, without looking at a remote version.

### Option 3: Multi-Module Layout (With go.work) without “common”

Similar to Option 2, but the “common” directory is not there. The common code is directly in the top-level “pkg”, “internal”, “tests”, ...

```
Unset
monorepo/
├─ pkg/           # The common source-code.
├─ tests/         # Tests for common source-code.
├─ types/         # Shared structs, interfaces, enums, constants, for
common.
├─ image/         # Independent Go module
│   └─ ...
│       └─ go.mod
├─ storage/       # Independent Go module
│   └─ ...
│       └─ go.mod
└─ go.work        # Ties them together
```

### Option 4: Bazel based monorepo

About Bazel:

- Open-source build and test tool similar to Make, Maven, and Gradle.
- Supports multiple languages and builds outputs for multiple platforms.
- Tracks dependencies between builds to rebuild only files which must be rebuilt.
- Supports remote caching.
- Runs only unit-tests which might be affected by the code change.

```

Unset
monorepo/
├── WORKSPACE                # Top-level file defining global bazel rules
├── BUILD.bazel              # (optional, often empty at root)
├── common/
│   ├── BUILD.bazel         # Build file defining go_library, go_test, ...
│   ├── go.mod
│   └── ...
├── storage/
│   ├── BUILD.bazel         # Build file defining go_library, go_test, ...
│   ├── go.mod
│   └── ...
├── image/
│   ├── BUILD.bazel         # Build file defining go_library, go_test, ...
│   ├── go.mod
│   └── ...

```

Alternatively, we can also move “common” directly to “pkg”, “internal” in the top-level monorepo directory. Similar to Option 3.

#### Pros:

- Similar to go.work, but much more advanced with more options for the future.
- Possibility to improve build time by smart caching.
  - This might be very important for testing PRs in case we include more repos in the monorepo in the future.
- Possibility to improve test time by running only tests affected by the code change in PR.
- Is a natural extension of go.work way - we can start with go.work and later replace it with Bazel without changing the project structure.
- We can also onboard our non-go projects to Bazel and build it using the same system - for example the aardvark-dns written in Rust.

#### Cons:

- The initial setup complexity is higher than with go.work.
  - But we are not the first one doing that: [here](#), [here](#), or [here](#).
- The learning curve is steeper for go developers. The “go build” or “go test” commands will not work.
  - But we are using “make” anyway, which can wrap bazel easily for the most common cases?
  - For developers who are new to go completely, it is not a big issue.
- Higher maintenance overhead.
  - The Build files must be updated from time to time when new dependencies are introduced. This can be somehow automated by using the [gazelle](#) tool.

## Generating the monorepo

[The git-filter-repo](#) is a python project allowing us to rewrite the git repository history. It can be used to move the whole repository to a single directory inside of another repository in a way that it includes the original git history.

In case we would need to move some files around in that repository, we can later use `git mv` to move them to the desired location.

The particular way to do that would depend on the Monorepo structure we choose.

## Issues workflow changes

There are currently following number of issues in the repositories in scope:

- Common - 49 issues
- Image - 93 issues
- Storage - 81 issues

When merging these three repositories, we would end up with 223 issues. We still want to be able to group the issues to particular project in the monorepo.